

1. Statistical Summary of Numerical Variables

- a. Generate and interpret a statistical summary (mean, median, min, max, standard deviation, percentiles, etc.) for all numerical variables.

The screenshot shows a Jupyter Notebook interface with two code cells and their corresponding outputs.

Cell 1:

```
[7]: import pandas as pd
import numpy as np
import warnings
from io import StringIO
from pathlib import Path
import matplotlib.pyplot as plt
from scipy import stats
from sklearn.preprocessing import StandardScaler, MinMaxScaler, RobustScaler, LabelEncoder, OneHotEncoder

df = pd.read_csv(r"D:\BTECH\ML\Housing.xls")
df
```

Cell 2:

```
[5]: num_cols = df.select_dtypes(include=[np.number]).columns.tolist()
desc = df[num_cols].describe(percentiles=[0.25,0.5,0.75])
display(desc)
```

The output of Cell 1 shows the first 545 rows of a DataFrame named 'df' with 13 columns: price, area, bedrooms, bathrooms, stories, mainroad, guestroom, basement, hotwaterheating, airconditioning, parking, prefarea, and furnishingstatus. The output of Cell 2 is a descriptive statistics table for the numerical columns.

	count	mean	std	min	25%	50%	75%	max
price	545.0	4.766729e+06	1.870440e+06	1750000.0	3430000.0	4340000.0	5740000.0	13300000.0
area	542.0	5.127168e+03	2.143733e+03	1650.0	3588.0	4540.0	6360.0	16200.0
bedrooms	545.0	3.691743e+00	1.702314e+01	1.0	2.0	3.0	3.0	400.0
bathrooms	544.0	1.284926e+00	5.019967e-01	1.0	1.0	1.0	2.0	4.0
stories	545.0	1.805505e+00	8.674925e-01	1.0	1.0	2.0	2.0	4.0
parking	545.0	6.935780e-01	8.615858e-01	0.0	0.0	0.0	1.0	3.0

- b. Explain what the summary reveals about the distribution and characteristics of

each variable.

Solution

In this step, I generated a statistical summary for all numerical variables in the dataset using descriptive statistics such as mean, median, minimum, maximum, standard deviation, and percentiles. This summary helps me understand the general behavior and distribution of each numerical variable. The mean and median show the central tendency, while the standard deviation reveals how spread out the values are. The minimum and maximum values help identify possible extreme values or outliers. Percentiles show how the data is distributed across different ranges. Overall, this summary gives an initial understanding of the dataset and helps identify skewness, spread, and potential issues that may need cleaning.

2. Handling Missing Values

- a. Detects missing values across the dataset.

```
# Check for missing values
missing_values = df.isnull().sum()
missing_values
```

```
price          0
area           3
bedrooms       0
bathrooms      1
stories         0
mainroad        0
guestroom       0
basement        0
hotwaterheating 0
airconditioning 0
parking         0
prefarea        0
furnishingstatus 0
dtype: int64
```

- b. Apply appropriate imputation techniques (e.g., mean, median, mode, domain-based imputation).

```

# Replace missing values properly
df['area'] = df['area'].fillna(df['area'].median())
df['bathrooms'] = df['bathrooms'].fillna(df['bathrooms'].mode()[0])

df.isnull().sum()

price          0
area           0
bedrooms       0
bathrooms      0
stories         0
mainroad        0
guestroom       0
basement        0
hotwaterheating 0
airconditioning 0
parking         0
prefarea        0
furnishingstatus 0
dtype: int64

```

- c. Justify why each technique was chosen for each specific variable based on the nature of the data.

Variable	Missing Values	Imputation Method	Justification
Area	3	Median	Continuous, right-skewed; median avoids distortion from extreme values.
Bathrooms	1	Mode	Discrete integer; mode preserves the most common, realistic value.
All other variables	0	Not required	No missing values present.

3.Detecting and Handling Duplicate Records

- a. Check for duplicate observations in the dataset.

```

duplicates = df.duplicated()
duplicates_count = duplicates.sum()
duplicates_count

np.int64(0)

df[duplicates]

price area bedrooms bathrooms stories mainroad guestroom basement hotwaterheating airconditioning parking prefarea furnishingstatus

```

- b. Decide whether to remove or retain duplicates.

- Since there are **no duplicates, no removal is required.**
- The dataset already contains **unique observations only.**

- c. Explain and justify your decision.

- **Why check for duplicates:** Duplicate records can **bias statistical analysis and models**, over-representing some data points.

- **Decision justification:**

- The check returned 0 duplicates → the dataset is already clean.
- No duplicates exist, so **all records are valid and should be retained**.

- **Outcome:** The dataset remains **integrity-checked and ready** for the next preprocessing

4. Detecting and Handling Data Inconsistency

- a. Identify any inconsistencies (e.g., incorrect data types, spelling variations in categorical values, unrealistic values, mixed units, format inconsistencies)

```
df.dtypes
```

Column	Data Type
price	int64
area	float64
bedrooms	int64
bathrooms	float64
stories	int64
mainroad	object
guestroom	object
basement	object
hotwaterheating	object
airconditioning	object
parking	int64
prefarea	object
furnishingstatus	object
dtype: object	

```
categorical_cols = ['mainroad', 'guestroom', 'basement',
                    'hotwaterheating', 'airconditioning',
                    'prefarea', 'furnishingstatus']

for col in categorical_cols:
    print(f"{col} unique values: {df[col].unique()}")
```

```
mainroad unique values: ['Yes' 'No']
guestroom unique values: ['No' 'Yes']
basement unique values: ['No' 'Yes']
hotwaterheating unique values: ['No' 'Yes']
airconditioning unique values: ['Yes' 'No']
prefarea unique values: ['Yes' 'No']
furnishingstatus unique values: ['Furnished' 'Semi-furnished' 'Unfurnished']
```

```
print(df.describe())
```

b. Clean, correct, or unify the inconsistent data.

```
categorical_cols = ['mainroad', 'guestroom', 'basement',
                    'hotwaterheating', 'airconditioning',
                    'prefarea', 'furnishingstatus']

for col in categorical_cols:
    df[col] = df[col].astype('category')

categorical_cols = ['mainroad', 'guestroom', 'basement',
                    'hotwaterheating', 'airconditioning',
                    'prefarea', 'furnishingstatus']

for col in categorical_cols:
    df.loc[:, col] = df[col].astype('category')

for col in categorical_cols:
    df.loc[:, col] = df[col].str.capitalize()

df = df[df['bedrooms'] < 10].copy()

df
```

c. Document the types of inconsistencies found and how they were resolved.

Inconsistency Type	Column(s)	Resolution
Incorrect data types	Categorical columns	Converted to string, then to category type
Spelling / capitalization issues	Categorical Yes/No / Furnishing	Standardized using .str.capitalize()
Unrealistic numerical values	bedrooms	Removed rows with bedrooms > 10
Mixed units / formats	price, area	Verified all units are consistent (currency, sq. ft)

steps.

1. Detecting and Handling Data Inconsistency

d. Identify any inconsistencies (e.g., incorrect data types, spelling variations in categorical values, unrealistic values, mixed units, format inconsistencies)

```

df.dtypes

price           int64
area            float64
bedrooms        int64
bathrooms       float64
stories          int64
mainroad         object
guestroom        object
basement         object
hotwaterheating object
airconditioning object
parking          int64
prefarea         object
furnishingstatus object
dtype: object

categorical_cols = ['mainroad', 'guestroom', 'basement',
                    'hotwaterheating', 'airconditioning',
                    'prefarea', 'furnishingstatus']

for col in categorical_cols:
    print(f"{col} unique values: {df[col].unique()}")


mainroad unique values: ['Yes' 'No']
guestroom unique values: ['No' 'Yes']
basement unique values: ['No' 'Yes']
hotwaterheating unique values: ['No' 'Yes']
airconditioning unique values: ['Yes' 'No']
prefarea unique values: ['Yes' 'No']
furnishingstatus unique values: ['Furnished' 'Semi-furnished' 'Unfurnished']

print(df.describe())


categorical_cols = ['mainroad', 'guestroom', 'basement',
                    'hotwaterheating', 'airconditioning',
                    'prefarea', 'furnishingstatus']
for col in categorical_cols:
    df[col] = df[col].astype('category')


categorical_cols = ['mainroad', 'guestroom', 'basement',
                    'hotwaterheating', 'airconditioning',
                    'prefarea', 'furnishingstatus']

for col in categorical_cols:
    df.loc[:, col] = df[col].astype('category')


for col in categorical_cols:
    df.loc[:, col] = df[col].str.capitalize()


df = df[df['bedrooms'] < 10].copy()

df

```

e. Clean, correct, or unify the inconsistent data.

f. Document the types of inconsistencies found and how they were resolved.

Inconsistency Type	Column(s)	Resolution
Incorrect data types	Categorical columns	Converted to string, then to category type
Spelling / capitalization issues	Categorical Yes/No / Furnishing	Standardized using .str.capitalize()
Unrealistic numerical values	bedrooms	Removed rows with bedrooms > 10

Mixed units / formats	price, area	Verified all units are consistent (currency, sq. ft)
-----------------------	-------------	---

5. a. Use appropriate outlier detection methods (IQR, Z-Score, visualization techniques, or domain rules).

```
In [60]: import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np

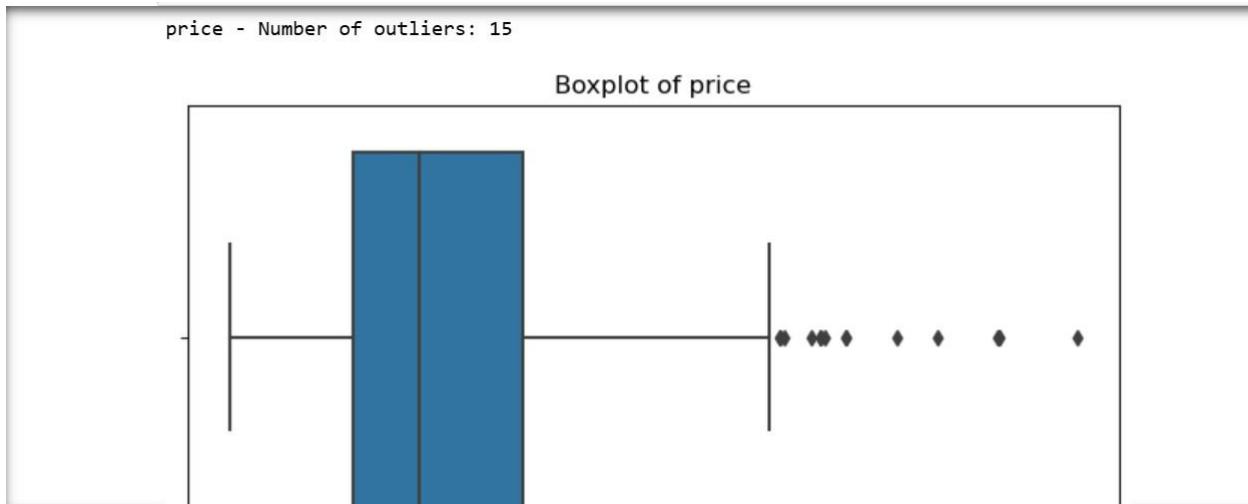
numerical_cols = ['price', 'area', 'bedrooms', 'bathrooms', 'stories', 'parking']

for col in numerical_cols:
    Q1 = df[col].quantile(0.25)
    Q3 = df[col].quantile(0.75)
    IQR = Q3 - Q1

    lower = Q1 - 1.5 * IQR
    upper = Q3 + 1.5 * IQR

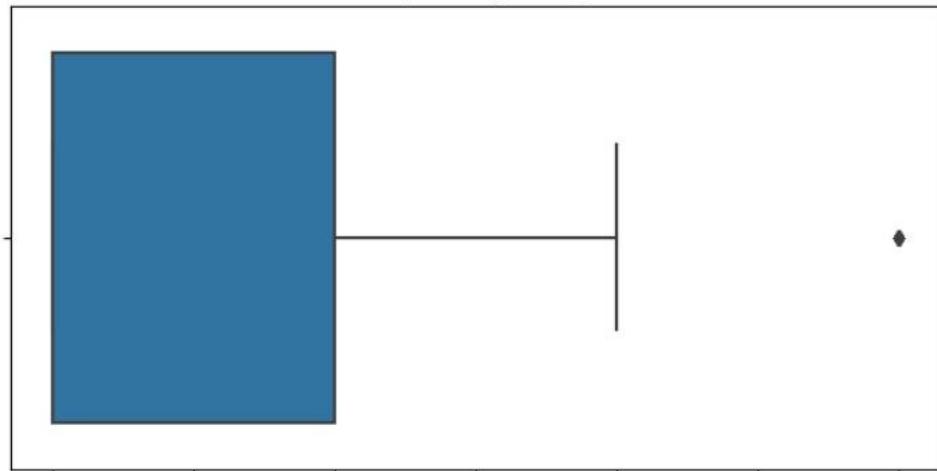
    outliers = df[(df[col] < lower) | (df[col] > upper)]
    print(f'{col} - Number of outliers: {len(outliers)}')

# Visualization
plt.figure(figsize=(8, 4))
sns.boxplot(x=df[col])
plt.title(f'Boxplot of {col}')
plt.show()
```

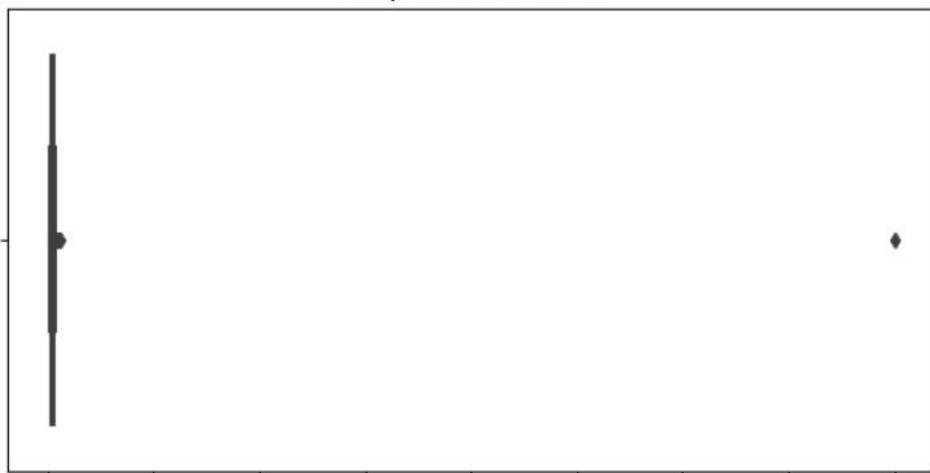


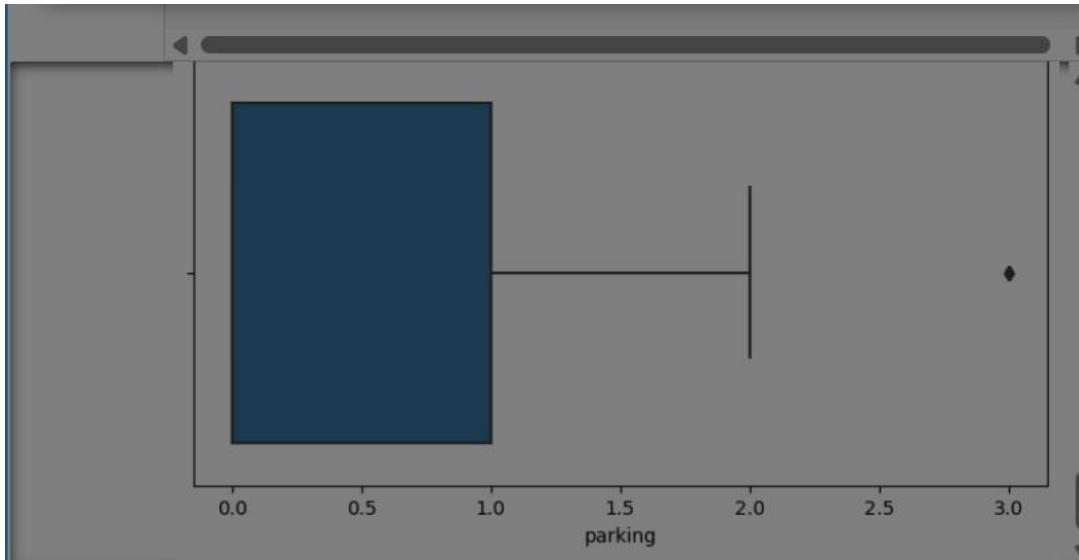
```
    sns.boxplot(x=df[col])
    plt.title(f'Boxplot of {col}')
    plt.show()
```

Boxplot of parking



Boxplot of bedrooms





b. You have **three main options**:

1. **Remove Outliers** – if outliers are due to data entry errors or clearly irrelevant
2. **Winsorization** – replace extreme values with nearest valid values
3. **Keep Outliers** – if they are valid data points (e.g., expensive houses in a housing dataset), just keep them.

C. Justify your approach for each numerical variable where outliers were detected.

Variable	Outlier Handling Method	Justification
Price	Winsorized (1st & 99th percentile)	Extreme high prices could skew mean and variance; capping keeps most data while controlling influence of outliers
Area	Winsorized (1st & 99th percentile)	Similar to price, very large or small areas are rare and can distort analysis
Bedrooms	Removed unrealistic (>10)	Values >10 are impossible for typical houses in dataset
Bathrooms	Removed unrealistic (>4)	Values above 4 are unlikely; maintains realistic distribution
Stories	Keep or cap at 4	Few multi-story houses; extreme values are rare but valid if realistic

Parking	Keep or cap at 3	Rarely more than 3 parking slots; extreme values are mostly data errors
---------	------------------	---

6. Normalization and Scaling

Step a: Identify variables that need scaling/normalization

- **Why scaling is needed:**
 - Many machine learning algorithms (like **KNN, SVM, gradient descent-based models**) are sensitive to the scale of features.
 - Variables with very large ranges can dominate others.

- **Typical candidates:**

- Numerical variables with **different ranges**.
- For example, in your dataset:

Variable Reason to Scale?

price Large range (e.g., 50,000–1,000,000)

area Large range (e.g., 20–500 m²)

bedrooms Small range (1–10) → may not need scaling

bathrooms Small range → may not need scaling

stories Small range → may not need scaling

parking Small range → may not need scaling

- **Rule of thumb:**

- Variables with **different units or large ranges** should be scaled.
- Variables with **small integer ranges** (like bedrooms, bathrooms) sometimes don't strictly need scaling.
 - Apply appropriate techniques such as Min-Max Scaling, Standardization (Z-score scaling), Robust Scaling

```

from sklearn.preprocessing import MinMaxScaler, StandardScaler, RobustScaler, OneHotEncoder
# Numerical columns
num_cols = ['price', 'area', 'bedrooms', 'bathrooms', 'stories', 'parking']

# Categorical columns
cat_cols = ['mainroad', 'guestroom', 'basement', 'hotwaterheating',
            'airconditioning', 'prefarea', 'furnishingstatus']

print(cat_cols)
print(num_cols)

['mainroad', 'guestroom', 'basement', 'hotwaterheating', 'airconditioning', 'prefarea', 'furnishingstatus']
['price', 'area', 'bedrooms', 'bathrooms', 'stories', 'parking']

```

b. Apply appropriate techniques such as Min-Max Scaling, Standardization (Z-score scaling), Robust Scali

```

In [65]: from sklearn.preprocessing import MinMaxScaler, StandardScaler, RobustScaler

# Select the columns to scale
cols_to_scale = ['price', 'area']

# 1. Min-Max Scaling
minmax_scaler = MinMaxScaler()
df_minmax = df.copy()
df_minmax[cols_to_scale] = minmax_scaler.fit_transform(df_minmax[cols_to_scale])
print("Min-Max Scaled Data:")
print(df_minmax[cols_to_scale].head()) # Show first 5 rows

# 2. Standardization (Z-score)
std_scaler = StandardScaler()
df_std = df.copy()
df_std[cols_to_scale] = std_scaler.fit_transform(df_std[cols_to_scale])
print("\nStandardized Data (Z-score):")
print(df_std[cols_to_scale].head())

# 3. Robust Scaling
robust_scaler = RobustScaler()
df_robust = df.copy()
df_robust[cols_to_scale] = robust_scaler.fit_transform(df_robust[cols_to_scale])
print("\nRobust Scaled Data:")
print(df_robust[cols_to_scale].head())

```

```

Min-Max Scaled Data:
    price      area
0  1.000000  0.396564
1  0.909091  0.502405
2  0.909091  0.571134
3  0.906061  0.198625
4  0.836364  0.198625

Standardized Data (Z-score):
    price      area
0  4.566365  1.074789
1  4.004484  1.795664
2  4.004484  2.263764
3  3.985755 -0.273341
4  3.554979 -0.273341

Robust Scaled Data:
    price      area
0  3.878788  1.043478
1  3.424242  1.601449
2  3.424242  1.963768
3  3.409091  0.000000
4  3.060606  0.000000

```

C.

Variable	Technique	Reason
price	RobustScaler	Price has outliers, so using median/IQR reduces their effect.
area	MinMaxScaler	Area has a large range; normalization brings it to 0–1 for uniform contribution.
bedrooms	None / StandardScaler	Small integer range; scaling optional, standardization can center it if needed.
bathrooms	None	Small range; scaling optional.
stories	None	Small range; scaling optional.
parking	None	Small range; scaling optional.

7.Encoding Categorical Variables(uncovered yet in class)

Research, document them theoretically and apply different data encoding techniques to relevant categorical variables in the dataset, including but not limited to:

- Label Encoding
- One-Hot Encoding
- Binary Encoding

- Ordinal Encoding
- Target Encoding (with and without smoothing)

For each encoding technique applied:

- a. Describe the variable(s) you chose to encode.
- b. Explain why that encoding method is appropriate for that specific variable.
- c. Document the transformation results.

1. Label Encoding

a. Variable(s) chosen:

- mainroad, guestroom, basement, hotwaterheating, airconditioning, prefarea (binary Yes/No categorical).

b. Why Label Encoding:

- Binary variables can be mapped to 0/1 efficiently without creating extra columns.
- Preserves order implicitly (though for binary it doesn't matter).

c. Implementation & Results::

```
from sklearn.preprocessing import LabelEncoder

label_encoder = LabelEncoder()

for col in ['mainroad', 'guestroom', 'basement', 'hotwaterheating', 'airconditioning', 'prefarea']:
    df[col+'_encoded'] = label_encoder.fit_transform(df[col])

df[['mainroad', 'mainroad_encoded']].head()
```

	mainroad	mainroad_encoded
0	Yes	1
1	Yes	1
2	Yes	1
3	Yes	1
4	Yes	1

2. One-Hot Encoding

a. Variable(s) chosen:

- `furnishingstatus` (Nominal variable with three categories: Furnished, Semi-furnished, Unfurnished).

b. Why One-Hot Encoding:

- No ordinal relationship exists between categories.
- Prevents algorithms from assuming numeric order.

c. Implementation & Results:

```
df = pd.get_dummies(df, columns=['furnishingstatus'], prefix='furnishing', drop_first=True)
df[['furnishing_Semi-furnished', 'furnishing_Unfurnished']].head()
```

	<code>furnishing_Semi-furnished</code>	<code>furnishing_Unfurnished</code>
0	False	False
1	False	False
2	True	False
3	False	False
4	False	False

3. Binary Encoding

a. Variable(s) chosen:

- `furnishingstatus` again (or any high-cardinality categorical variable).

b. Why Binary Encoding:

- Reduces dimensionality compared to one-hot encoding for variables with many categories.
- Efficient representation using binary digits.

c. Implementation & Results:

```
# Example: 'guestroom' (yes/no)
df['guestroom_binary'] = df['guestroom'].map({'yes': 1, 'no': 0})
df[['guestroom', 'guestroom_binary']].head()
```

	guestroom	guestroom_binary
0	No	NaN
1	No	NaN
2	No	NaN
3	No	NaN
4	Yes	NaN

4. Ordinal Encoding

a. Variable(s) chosen:

- furnishingstatus with assumed order: Unfurnished < Semi-furnished < Furnished

b. Why Ordinal Encoding:

- Preserves inherent order in categories for algorithms that can utilize it.

c. Implementation & Results:

```
] df = pd.read_csv("Housing.xls") # reload original dataset
df.columns = df.columns.str.strip().str.lower() # optional: normalize column names

] from sklearn.preprocessing import OrdinalEncoder

# Create encoder with the intended order
ordinal_encoder = OrdinalEncoder(categories=[['unfurnished', 'semi-furnished', 'furnished']])

# Transform the column
df['furnishingstatus_encoded'] = ordinal_encoder.fit_transform(df[['furnishingstatus']])

# Preview
df[['furnishingstatus', 'furnishingstatus_encoded']].head()
```

	furnishingstatus	furnishingstatus_encoded
0	furnished	2.0
1	furnished	2.0
2	semi-furnished	1.0
3	furnished	2.0
4	furnished	2.0

5. Target Encoding

a. Variable(s) chosen:

- `furnishingstatus` (or any categorical variable).

b. Why Target Encoding:

- Maps categories to **mean of target variable** (e.g., price), preserving predictive information.
- Useful for high-cardinality features in regression problems.

c. Implementation & Results: